

Question 2

```
In [1]: from tensorflow import keras
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from ultralytics import YOLO
import cv2
import pickle
from pathlib import Path
import matplotlib.pyplot as plt

import numpy as np
import os
```

2024-12-04 23:43:07.866875: I tensorflow/core/platform/cpu_feature_guard.cc:210] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.
To enable the following instructions: SSE4.1 SSE4.2 AVX AVX2 FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.

```
In [3]: import os
import cv2
import torch

# Paths
model_path = './runs/detect/car_detection_20/weights/best.pt'
input_dir = 'car_detection_dataset/testing_images/'
output_dir = 'output_images/'

# Create output directory
os.makedirs(output_dir, exist_ok=True)

# Limit GPU memory usage and clear cache
torch.cuda.set_per_process_memory_fraction(0.8, device=0)
torch.cuda.empty_cache()

# Load YOLO model
model = YOLO(model_path)

def resize_image(image_path, max_size=640):
    """
    Resize the image to a maximum size while maintaining aspect ratio.
    """
    image = cv2.imread(image_path)
    h, w = image.shape[:2]
    scale = max_size / max(h, w)
    new_h, new_w = int(h * scale), int(w * scale)
    resized_image = cv2.resize(image, (new_w, new_h))
    return resized_image, scale

def draw_boxes_and_labels(img, result, scale=1):
    """
    Draw bounding boxes and labels on the image.
    Scale is used to adjust bounding box coordinates if the image was resized.
    """
    for box, conf, cls in zip(result.bboxes.xyxy, result.bboxes.conf, result.bboxes.cls):
        # Scale bounding box coordinates back to the resized image
        x1, y1, x2, y2 = (box.cpu().numpy() * scale).astype(int)

        # Draw bounding box
        cv2.rectangle(img, (x1, y1), (x2, y2), (0, 255, 0), 2)

        # Add label with confidence score
        label = f'Car {conf:.2f}'
        cv2.putText(img, label, (x1, y1 - 10), cv2.FONT_HERSHEY_SIMPLEX, 0.9, (0, 255, 0), 2)
    return img

# Process each image in the input directory
input_files = [os.path.join(input_dir, f) for f in os.listdir(input_dir) if f.lower().endswith(('.png', '.jpg', '.jpeg'))]
```

```
for file_path in input_files:
    print(f"Processing: {file_path}")

    # Resize image
    resized_img, scale = resize_image(file_path)

    # Get predictions from the model
    results = model(file_path, verbose=False)

    # Draw bounding boxes and labels on the resized image
    for r in results:
        resized_img = draw_boxes_and_labels(resized_img, r, scale)

    # Save the processed image to the output directory
    output_path = os.path.join(output_dir, os.path.basename(file_path))
    cv2.imwrite(output_path, resized_img)

print("All images processed and saved ")
```

[illegible]

[illegible]

```

Processing: car_detection_dataset/testing_images/vid_5_31100.jpg
Processing: car_detection_dataset/testing_images/vid_5_30820.jpg
Processing: car_detection_dataset/testing_images/vid_5_29500.jpg
Processing: car_detection_dataset/testing_images/vid_5_27760.jpg
Processing: car_detection_dataset/testing_images/vid_5_29840.jpg
Processing: car_detection_dataset/testing_images/vid_5_26860.jpg
Processing: car_detection_dataset/testing_images/vid_5_29820.jpg
Processing: car_detection_dataset/testing_images/vid_5_30700.jpg
Processing: car_detection_dataset/testing_images/vid_5_27420.jpg
Processing: car_detection_dataset/testing_images/vid_5_28180.jpg
Processing: car_detection_dataset/testing_images/vid_5_30000.jpg
Processing: car_detection_dataset/testing_images/vid_5_29520.jpg
Processing: car_detection_dataset/testing_images/vid_5_28700.jpg
Processing: car_detection_dataset/testing_images/vid_5_27300.jpg
Processing: car_detection_dataset/testing_images/vid_5_31480.jpg
Processing: car_detection_dataset/testing_images/vid_5_28540.jpg
Processing: car_detection_dataset/testing_images/vid_5_27720.jpg
Processing: car_detection_dataset/testing_images/vid_5_30140.jpg
Processing: car_detection_dataset/testing_images/vid_5_31180.jpg
Processing: car_detection_dataset/testing_images/vid_5_27920.jpg
Processing: car_detection_dataset/testing_images/vid_5_27980.jpg
Processing: car_detection_dataset/testing_images/vid_5_27460.jpg
Processing: car_detection_dataset/testing_images/vid_5_27780.jpg
Processing: car_detection_dataset/testing_images/vid_5_25100.jpg
Processing: car_detection_dataset/testing_images/vid_5_30040.jpg
Processing: car_detection_dataset/testing_images/vid_5_27560.jpg
Processing: car_detection_dataset/testing_images/vid_5_31380.jpg
Processing: car_detection_dataset/testing_images/vid_5_27680.jpg
Processing: car_detection_dataset/testing_images/vid_5_27320.jpg
Processing: car_detection_dataset/testing_images/vid_5_30760.jpg
Processing: car_detection_dataset/testing_images/vid_5_26400.jpg
Processing: car_detection_dataset/testing_images/vid_5_26820.jpg
Processing: car_detection_dataset/testing_images/vid_5_27860.jpg
Processing: car_detection_dataset/testing_images/vid_5_27360.jpg
Processing: car_detection_dataset/testing_images/vid_5_27520.jpg
All images processed and saved

```

Question 3

```

In [4]: import os
from pathlib import Path
import cv2
import matplotlib.pyplot as plt
from ultralytics import YOLO

def main():
    # Load YOLO model
    modelPath = './runs/detect/car_detection_20/weights/best.pt'
    imageDir = './Make sense/'
    labelDir = './Make sense labels/'
    model = YOLO(modelPath)

    # Get List of image files
    imageFiles = [file for file in os.listdir(imageDir) if file.lower().endswith(('.png', '.jpg', '.jpeg'))]

    totalCarsAll = 0
    detectedCarsAll = 0

    for imageFile in imageFiles:
        # Paths for image and label files
        imagePath = Path(imageDir) / imageFile
        labelPath = Path(labelDir) / f"{Path(imageFile).stem}.txt"

        # Load image
        image = cv2.imread(str(imagePath))
        originalImage = image.copy()
        trueBoxesImage = image.copy()
        bothBoxesImage = image.copy()

```

```

# Load true boxes
trueBoxes = []
if labelPath.is_file():
    with open(labelPath, 'r') as lp:
        for line in lp:
            classId, x, y, w, h = map(float, line.strip().split())
            imageHeight, imageWidth = image.shape[:2]
            x1, y1 = int((x - w / 2) * imageWidth), int((y - h / 2) * imageHeight)
            x2, y2 = int((x + w / 2) * imageWidth), int((y + h / 2) * imageHeight)
            trueBoxes.append([x1, y1, x2, y2])
            cv2.rectangle(trueBoxesImage, (x1, y1), (x2, y2), (255, 0, 0), 2)

# Get predictions
results = model(image)
predictedBoxes = results[0].boxes.xyxy.cpu().numpy()

# Draw predicted boxes
for box in predictedBoxes:
    x1, y1, x2, y2 = map(int, box[:4])
    cv2.rectangle(bothBoxesImage, (x1, y1), (x2, y2), (0, 255, 0), 2)

# Draw true boxes on combined image
for box in trueBoxes:
    x1, y1, x2, y2 = box
    cv2.rectangle(bothBoxesImage, (x1, y1), (x2, y2), (255, 0, 0), 2)

# Plot results
plt.figure(figsize=(15, 5))
plt.subplot(131)
plt.imshow(cv2.cvtColor(originalImage, cv2.COLOR_BGR2RGB))
plt.title('Original Image')
plt.axis('off')

plt.subplot(132)
plt.imshow(cv2.cvtColor(trueBoxesImage, cv2.COLOR_BGR2RGB))
plt.title('True Boxes (Blue)')
plt.axis('off')

plt.subplot(133)
plt.imshow(cv2.cvtColor(bothBoxesImage, cv2.COLOR_BGR2RGB))
plt.title('Predicted (Green) & True (Blue) Boxes')
plt.axis('off')

plt.tight_layout()
plt.show()

# Evaluate predictions directly
carsIdentified = 0
totalCars = len(trueBoxes)
matchedPreds = set()

for trueBox in trueBoxes:
    bestIoU = 0.5
    bestPredIdx = -1

    for j, predBox in enumerate(predictedBoxes):
        if j in matchedPreds:
            continue

        # IoU calculation logic
        b1 = predBox
        b2 = [trueBox[0] - 5, trueBox[1] - 5,
              trueBox[2] + 5, trueBox[3] + 5]

        interX1 = max(b1[0], b2[0])
        interY1 = max(b1[1], b2[1])
        interX2 = min(b1[2], b2[2])
        interY2 = min(b1[3], b2[3])

        interArea = max(0, interX2 - interX1) * max(0, interY2 - interY1)

```

```

        predArea = (b1[2] - b1[0]) * (b1[3] - b1[1])
        trueArea = (b2[2] - b2[0]) * (b2[3] - b2[1])

        iou = interArea / float(predArea + trueArea - interArea)

        if iou > bestIoU:
            bestIoU = iou
            bestPredIdx = j

    if bestPredIdx >= 0:
        carsIdentified += 1
        matchedPreds.add(bestPredIdx)

# Accumulate results
totalCarsAll += totalCars
detectedCarsAll += carsIdentified

print(f"\nImage: {imageFile}")
print(f"Cars in image: {totalCars}")
print(f"Cars detected: {carsIdentified}")

# Final results
print(f"\nTotal Results:")
print(f"Total cars in all images: {totalCarsAll}")
print(f"Total cars detected: {detectedCarsAll}")
print(f"Detection accuracy: {detectedCarsAll / totalCarsAll:.2f}")

if __name__ == "__main__":
    main()

```

0: 384x640 1 car, 7.8ms

Speed: 1.8ms preprocess, 7.8ms inference, 1.0ms postprocess per image at shape (1, 3, 384, 640)




Image: vid_5_27860.jpg

Cars in image: 1

Cars detected: 1

0: 384x640 2 cars, 7.6ms

Speed: 1.3ms preprocess, 7.6ms inference, 1.0ms postprocess per image at shape (1, 3, 384, 640)

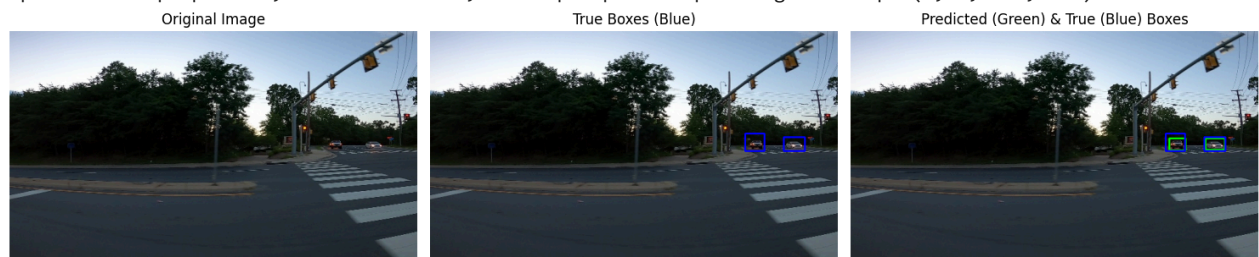


Image: vid_5_27320.jpg

Cars in image: 2

Cars detected: 0

0: 384x640 1 car, 7.6ms

Speed: 1.2ms preprocess, 7.6ms inference, 1.0ms postprocess per image at shape (1, 3, 384, 640)



Image: vid_5_27240.jpg
 Cars in image: 1
 Cars detected: 1

0: 384x640 (no detections), 7.6ms
 Speed: 1.3ms preprocess, 7.6ms inference, 0.4ms postprocess per image at shape (1, 3, 384, 640)




Image: vid_5_27720.jpg
 Cars in image: 0
 Cars detected: 0

0: 384x640 (no detections), 7.5ms
 Speed: 1.3ms preprocess, 7.5ms inference, 0.4ms postprocess per image at shape (1, 3, 384, 640)



Image: vid_5_26980.jpg
 Cars in image: 0
 Cars detected: 0

Total Results:
 Total cars in all images: 4
 Total cars detected: 2
 Detection accuracy: 0.50

In []: