

```
In [2]: import numpy as np
import numpy.random as npr
import matplotlib.pyplot as plt

import tensorflow as tf
from tensorflow import keras
from keras.wrappers.scikit_learn import KerasClassifier
import keras_tuner as kt
from tensorflow.keras import layers, models, optimizers
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, f1_score
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint

%matplotlib inline
plt.style.use('bmh')
```


```
In [3]: class_names = ['Roses', 'Magnolias', 'Lilies', 'Sunflowers', 'Orchids',
'Marigold', 'Hibiscus', 'Firebush', 'Pentas', 'Bougainvillea']
X_train = np.load('flower_species_classification/data_train.npy').T
t_train = np.load('flower_species_classification/labels_train.npy')

print(X_train.shape, t_train.shape)
```

(1658, 270000) (1658,)

```
In [4]: vals, counts = np.unique(t_train, return_counts=True)

plt.bar(vals, counts)
plt.xticks(range(10), range(10))
plt.xlabel('Classes', size=20)
plt.ylabel('# Samples per Class', size=20)
plt.title('Training Data (Total = '+str(X_train.shape[1])+ ' samples)', size=15);
```



```
In [5]: X_train, X_val, t_train, t_val = train_test_split(X_train, t_train, test_size=0.2, stratify=t_train, shuffle
```

```
In [6]: X_train_rs = tf.constant(X_train.reshape((X_train.shape[0],300,300,3))/255, dtype=tf.float32)
X_val_rs = tf.constant(X_val.reshape((X_val.shape[0],300,300,3))/255, dtype=tf.float32)
```

2024-12-04 21:12:43.458957: I tensorflow/core/platform/cpu\_feature\_guard.cc:151] This TensorFlow binary is optimized with oneAPI Deep Neural Network Library (oneDNN) to use the following CPU instructions in performance-critical operations: SSE4.1 SSE4.2 AVX AVX2 FMA  
To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.  
2024-12-04 21:12:44.847951: I tensorflow/core/common\_runtime/gpu/gpu\_device.cc:1525] Created device /job:localhost/replica:0/task:0/device:GPU:0 with 78045 MB memory: -> device: 0, name: NVIDIA A100-SXM4-80GB, pci bus id: 0000:47:00.0, compute capability: 8.0  
2024-12-04 21:12:44.850844: I tensorflow/core/common\_runtime/gpu/gpu\_device.cc:1525] Created device /job:localhost/replica:0/task:0/device:GPU:1 with 78902 MB memory: -> device: 1, name: NVIDIA A100-SXM4-80GB, pci bus id: 0000:4e:00.0, compute capability: 8.0

```
In [7]: def build_model(hp):
model = models.Sequential()
```

```

# Single Convolutional Layer with hyperparameters
model.add(layers.Conv2D(
    filters=hp.Choice('filters', values=[32, 64]), # Corrected hyperparameter name
    kernel_size=hp.Choice('kernel_size', values=[3, 5]), # Corrected hyperparameter name
    activation='relu',
    input_shape=(300, 300, 3)
))
model.add(layers.MaxPooling2D((2, 2)))

model.add(layers.Flatten())

# Single Dense Layer with hyperparameters
model.add(layers.Dense(
    units=hp.Choice('dense_units', values=[128, 256]), # Corrected hyperparameter name
    activation='relu'
))

model.add(layers.Dense(10, activation='softmax')) # Output Layer with 10 classes

# Compile the model with a tunable learning rate
model.compile(
    optimizer=optimizers.Adam(
        learning_rate=hp.Choice('learning_rate', values=[1e-3, 1e-4]) # Corrected hyperparameter name
    ),
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy']
)

return model

```

```

In [8]: # Set up Keras Tuner for the simplified model
tuner = kt.RandomSearch(
    build_model,
    objective='val_accuracy',
    max_trials=10, # Number of hyperparameter configurations to try
    executions_per_trial=1,
    directory='flower_tuning',
    project_name='flower_species'
)

# Early stopping to prevent overfitting
stop_early = EarlyStopping(monitor='val_loss', patience=3)

# Perform the hyperparameter search
tuner.search(
    X_train_rs, t_train,
    epochs=10,
    validation_data=(X_val_rs, t_val),
    callbacks=[stop_early]
)

# Retrieve the best hyperparameters found
best_hps = tuner.get_best_hyperparameters(num_trials=1)[0]

# Print the best hyperparameters
print(f"""
filters for convolutional layer: {best_hps.get('filters')}
kernel size for convolutional layer: {best_hps.get('kernel_size')}
dense units: {best_hps.get('dense_units')}
learning rate: {best_hps.get('learning_rate')}
""")

```

```
INFO:tensorflow:Reloading Oracle from existing project flower_tuning/flower_species/oracle.json
INFO:tensorflow:Reloading Tuner from flower_tuning/flower_species/tuner0.json
INFO:tensorflow:Oracle triggered exit
```

```
filters for convolutional layer: 64
kernel size for convolutional layer: 3
dense units: 256
learning rate: 0.0001
```

```
In [9]: best_model = tuner.hypermodel.build(best_hps)

history = best_model.fit(
    X_train_rs, t_train,
    epochs=20,
    validation_data=(X_val_rs, t_val),
    callbacks=[stop_early]
)

train_loss, train_accuracy = best_model.evaluate(X_train_rs, t_train)
print(f"Training Accuracy: {train_accuracy:.2f}")

val_loss, val_accuracy = best_model.evaluate(X_val_rs, t_val)
print(f"Validation Accuracy: {val_accuracy:.2f}")
```

Epoch 1/20

```
2024-12-04 21:13:01.219473: I tensorflow/stream_executor/cuda/cuda_dnn.cc:366] Loaded cuDNN version 8201
2024-12-04 21:13:02.748268: I tensorflow/stream_executor/cuda/cuda_blas.cc:1774] TensorFloat-32 will be used
for the matrix multiplication. This will only be logged once.
```

```
42/42 [=====] - 5s 37ms/step - loss: 6.3572 - accuracy: 0.2081 - val_loss: 1.9455 -
val_accuracy: 0.3855
```

Epoch 2/20

```
42/42 [=====] - 1s 30ms/step - loss: 1.3396 - accuracy: 0.5385 - val_loss: 1.4422 -
val_accuracy: 0.4849
```

Epoch 3/20

```
42/42 [=====] - 1s 30ms/step - loss: 0.8771 - accuracy: 0.7278 - val_loss: 1.3237 -
val_accuracy: 0.5482
```

Epoch 4/20

```
42/42 [=====] - 1s 31ms/step - loss: 0.5230 - accuracy: 0.8861 - val_loss: 1.1657 -
val_accuracy: 0.5693
```

Epoch 5/20

```
42/42 [=====] - 1s 31ms/step - loss: 0.3320 - accuracy: 0.9321 - val_loss: 1.2516 -
val_accuracy: 0.5542
```

Epoch 6/20

```
42/42 [=====] - 1s 30ms/step - loss: 0.2973 - accuracy: 0.9336 - val_loss: 1.1585 -
val_accuracy: 0.6114
```

Epoch 7/20

```
42/42 [=====] - 1s 31ms/step - loss: 0.1517 - accuracy: 0.9849 - val_loss: 1.1971 -
val_accuracy: 0.6054
```

Epoch 8/20

```
42/42 [=====] - 1s 31ms/step - loss: 0.0908 - accuracy: 0.9947 - val_loss: 1.0321 -
val_accuracy: 0.6416
```

Epoch 9/20

```
42/42 [=====] - 1s 32ms/step - loss: 0.0663 - accuracy: 0.9977 - val_loss: 1.0621 -
val_accuracy: 0.6355
```

Epoch 10/20

```
42/42 [=====] - 1s 31ms/step - loss: 0.0476 - accuracy: 1.0000 - val_loss: 1.1397 -
val_accuracy: 0.6175
```

Epoch 11/20

```
42/42 [=====] - 1s 30ms/step - loss: 0.0380 - accuracy: 1.0000 - val_loss: 1.0801 -
val_accuracy: 0.6446
```

```
42/42 [=====] - 0s 11ms/step - loss: 0.0306 - accuracy: 1.0000
```

```
Training Accuracy: 1.00
```

```
11/11 [=====] - 0s 11ms/step - loss: 1.0801 - accuracy: 0.6446
```

```
Validation Accuracy: 0.64
```

```
In [10]: train_predictions = tf.argmax(best_model.predict(X_train_rs), axis=1).numpy()
val_predictions = tf.argmax(best_model.predict(X_val_rs), axis=1).numpy()

print("\nClassification Report on Train Data:")
```

```
print(classification_report(t_train, train_predictions, target_names=class_names))

print("\nClassification Report on Validation Data:")
print(classification_report(t_val, val_predictions, target_names=class_names))
```


Classification Report on Train Data:

	precision	recall	f1-score	support
Roses	1.00	1.00	1.00	141
Magnolias	1.00	1.00	1.00	144
Lilies	1.00	1.00	1.00	164
Sunflowers	1.00	1.00	1.00	112
Orchids	1.00	1.00	1.00	138
Marigold	1.00	1.00	1.00	125
Hibiscus	1.00	1.00	1.00	128
Firebush	1.00	1.00	1.00	138
Pentas	1.00	1.00	1.00	130
Bougainvillea	1.00	1.00	1.00	106
accuracy			1.00	1326
macro avg	1.00	1.00	1.00	1326
weighted avg	1.00	1.00	1.00	1326

Classification Report on Validation Data:

	precision	recall	f1-score	support
Roses	0.61	0.64	0.62	36
Magnolias	0.67	0.86	0.76	36
Lilies	0.41	0.32	0.36	41
Sunflowers	0.88	0.75	0.81	28
Orchids	0.70	0.74	0.72	35
Marigold	0.68	0.87	0.76	31
Hibiscus	0.58	0.59	0.58	32
Firebush	0.83	0.74	0.78	34
Pentas	0.74	0.44	0.55	32
Bougainvillea	0.45	0.56	0.50	27
accuracy			0.64	332
macro avg	0.65	0.65	0.64	332
weighted avg	0.65	0.64	0.64	332

```
In [11]: plt.figure(figsize=(12, 6))
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.legend()
plt.title('Learning Curves')
plt.xlabel('Epochs')
plt.ylabel('Metrics')
plt.show()
```



```
In [12]: best_model.save('best_model.keras', save_format='keras')
```

```
In [ ]:
```