

```
In [1]: import numpy as np
import tensorflow as tf
from tensorflow import keras
from keras.applications import VGG16
from keras.models import Sequential
from keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
from keras.callbacks import ModelCheckpoint, EarlyStopping
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelBinarizer
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from ultralytics import YOLO
import matplotlib.pyplot as plt
import cv2
import pandas as pd
from pathlib import Path
import yaml
import shutil
import pickle

import os
```

2024-12-04 23:55:39.977758: I tensorflow/core/platform/cpu_feature_guard.cc:210] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.
To enable the following instructions: SSE4.1 SSE4.2 AVX AVX2 FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.

Question 2

```
In [2]: df = pd.read_csv('./car_detection_dataset/train_bounding_boxes.csv')
image_files = sorted([l for l in os.listdir('./car_detection_dataset/training_images')
                      if l.lower().endswith(('.png', '.jpg', '.jpeg'))
                      and os.path.isfile(os.path.join('./car_detection_dataset/training_images', l))])

train_images, val_images = train_test_split(image_files, test_size=0.2, random_state=42)

os.makedirs('./car_detection_dataset/train/image', exist_ok=True)
os.makedirs('./car_detection_dataset/train/label', exist_ok=True)
os.makedirs('./car_detection_dataset/val/image', exist_ok=True)
os.makedirs('./car_detection_dataset/val/label', exist_ok=True)
```

```
In [3]: # Process and move training images
for images in train_images:
    source = os.path.join('./car_detection_dataset/training_images', images)

    # Copy image to the destination folder
    destination = os.path.join('./car_detection_dataset/train/images', images)
    shutil.copy(source, destination)

    # Retrieve bounding box data for the current image
    image_data = df[df['image'] == images]
    label_path = os.path.join('./car_detection_dataset/train/labels', f'{Path(images).stem}.txt')

    # Open the label file and write the bounding box annotations
    with open(label_path, 'w') as lp:
        if not image_data.empty:
            image = cv2.imread(source)
            for _, row in image_data.iterrows():
                # Normalize the bounding box coordinates
                x = (row['xmin'] + row['xmax']) / 2 / image.shape[1]
                y = (row['ymin'] + row['ymax']) / 2 / image.shape[0]
                w = (row['xmax'] - row['xmin']) / image.shape[1]
                h = (row['ymax'] - row['ymin']) / image.shape[0]

                # Write the annotation (label) in YOLO format
                lp.write(f'0 {x} {y} {w} {h}\n')
```

```
In [4]: # Process and move validation images
for images in val_images:
    source = os.path.join('./car_detection_dataset/training_images', images)
    # Copy image to the destination folder
    destination = os.path.join('./car_detection_dataset/val/images', images)
    shutil.copy(source, destination)

    # Retrieve bounding box data for the current image
    image_data = df[df['image'] == images]
    label_path = os.path.join('./car_detection_dataset/val/labels', f'{Path(images).stem}.txt')

    # Open the label file and write the bounding box annotations
    with open(label_path, 'w') as lp:
        if not image_data.empty:
            image = cv2.imread(source)
            for _, row in image_data.iterrows():
                # Normalize the bounding box coordinates
                x = (row['xmin'] + row['xmax']) / 2 / image.shape[1]
                y = (row['ymin'] + row['ymax']) / 2 / image.shape[0]
                w = (row['xmax'] - row['xmin']) / image.shape[1]
                h = (row['ymax'] - row['ymin']) / image.shape[0]

                # Write the annotation (label) in YOLO format
                lp.write(f'0 {x} {y} {w} {h}\n')
```

```
In [5]: data_yaml = {
    'train': './car_detection_dataset/train/images',
    'val': './car_detection_dataset/val/images',
    'nc': 1, # Number of classes (in this case, only 'car')
    'names': ['car']
}

# Save the YAML configuration to a file
import yaml
with open('data.yaml', 'w') as f:
    yaml.dump(data_yaml, f)
```

```
In [6]: maxWidth, maxHeight = 0, 0
for name in os.listdir('./car_detection_dataset/training_images/'):
    imagePath = os.path.join('./car_detection_dataset/training_images/', name)
    # Read the image
    image = cv2.imread(imagePath)
    if image is not None: # Ensure the image was read successfully
        height = image.shape[0]
        width = image.shape[1] # Get the dimensions
        # Update max width and height if necessary
        if width > maxWidth:
            maxWidth = width
        if height > maxHeight:
            maxHeight = height

# Output the maximum width and height
print(f"Max Width: {maxWidth}, Max Height: {maxHeight}")
```

Max Width: 676, Max Height: 380

```
In [10]: # Define file paths
pickle_file = 'model_info.pkl'
dataset_yaml = 'data.yaml'
results_csv = ""
model_path = ""

# Set batch sizes for training
best_map = 0
best_model_info = {}

# Load dataset configuration from YAML file
with open(dataset_yaml, 'r') as f:
    dataset_config = yaml.safe_load(f)
```

```

epoch_options = [10, 20]
# Iterate over batch sizes to train and select the best model
for epoch in epoch_options:
    results_csv = f'./runs/detect/car_detection_{epoch}/results.csv'
    model_path = f'./runs/detect/car_detection_{epoch}/weights/best.pt'

    if not os.path.exists(model_path): # If model hasn't been trained
        try:
            # Determine image size (based on dataset)
            max_width = max_height = 0
            image = max(32 * round(max(max_width, max_height) / 32), 640)

            # Initialize YOLOv5m model and start training
            model = YOLO('yolov5m.pt')
            model.train(
                data=dataset_yaml,
                epochs=epoch,
                imgsz=image,
                batch=16,
                name=f'car_detection_{epoch}'
            )
        except Exception as e:
            print(f"Error during training with batch size {epoch}: {e}")
            continue

    # Read the results CSV after training
    if os.path.exists(results_csv):
        df = pd.read_csv(results_csv)
        current_map = df['metrics/precision(B)'].iloc[-1]

        # Update best model if current mAP is higher
        if current_map > best_map:
            best_map = current_map
            best_model_info = {'model_path': model_path, 'results_csv': results_csv}
    else:
        print(f"Results file for batch {epoch} not found.")

```

PRO TIP 💡 Replace 'model=yolov5m.pt' with new 'model=yolov5mu.pt'.

YOLOv5 'u' models are trained with <https://github.com/ultralytics/ultralytics> and feature improved performance vs standard YOLOv5 models trained with <https://github.com/ultralytics/yolov5>.

New <https://pypi.org/project/ultralytics/8.3.41> available 😊 Update with 'pip install -U ultralytics'

Ultralytics 8.3.39 🚀 Python-3.12.7 torch-2.5.1+cu124 CUDA:0 (NVIDIA A100-SXM4-80GB, 81038MiB)

engine/trainer: task=detect, mode=train, model=yolov5m.pt, data=data.yaml, epochs=10, time=None, patience=100, batch=16, imgsz=640, save=True, save_period=-1, cache=False, device=None, workers=8, project=None, name=car_detection_10, exist_ok=False, pretrained=True, optimizer=auto, verbose=True, seed=0, deterministic=True, single_cls=False, rect=False, cos_lr=False, close_mosaic=10, resume=False, amp=True, fraction=1.0, profile=False, freeze=None, multi_scale=False, overlap_mask=True, mask_ratio=4, dropout=0.0, val=True, split=val, save_json=False, save_hybrid=False, conf=None, iou=0.7, max_det=300, half=False, dnn=False, plots=True, source=None, vid_stride=1, stream_buffer=False, visualize=False, augment=False, agnostic_nms=False, classes=None, retina_masks=False, embed=None, show=False, save_frames=False, save_txt=False, save_conf=False, save_crop=False, show_labels=True, show_conf=True, show_boxes=True, line_width=None, format=torchscript, keras=False, optimize=False, int8=False, dynamic=False, simplify=True, opset=None, workspace=None, nms=False, lr0=0.01, lrf=0.01, momentum=0.937, weight_decay=0.0005, warmup_epochs=3.0, warmup_momentum=0.8, warmup_bias_lr=0.1, box=7.5, cls=0.5, dfl=1.5, pose=12.0, kobj=1.0, nbs=64, hsv_h=0.015, hsv_s=0.7, hsv_v=0.4, degrees=0.0, translate=0.1, scale=0.5, shear=0.0, perspective=0.0, flipud=0.0, fliplr=0.5, bgr=0.0, mosaic=1.0, mixup=0.0, copy_paste=0.0, copy_paste_mode=flip, auto_augment=randaugment, erasing=0.4, crop_fraction=1.0, cfg=None, tracker=botsort.yaml, save_dir=runs/detect/car_detection_10

Overriding model.yaml nc=80 with nc=1

		from	n	params	module	arguments
0		-1	1	5280	ultralytics.nn.modules.conv.Conv	[3, 48, 6, 2, 2]
1		-1	1	41664	ultralytics.nn.modules.conv.Conv	[48, 96, 3, 2]
2		-1	2	65280	ultralytics.nn.modules.block.C3	[96, 96, 2]
3		-1	1	166272	ultralytics.nn.modules.conv.Conv	[96, 192, 3, 2]
4		-1	4	444672	ultralytics.nn.modules.block.C3	[192, 192, 4]
5		-1	1	664320	ultralytics.nn.modules.conv.Conv	[192, 384, 3, 2]
6		-1	6	2512896	ultralytics.nn.modules.block.C3	[384, 384, 6]
7		-1	1	2655744	ultralytics.nn.modules.conv.Conv	[384, 768, 3, 2]
8		-1	2	4134912	ultralytics.nn.modules.block.C3	[768, 768, 2]
9		-1	1	1476864	ultralytics.nn.modules.block.SPPF	[768, 768, 5]
10		-1	1	295680	ultralytics.nn.modules.conv.Conv	[768, 384, 1, 1]
11		-1	1	0	torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
12	[-1, 6]	1	1	0	ultralytics.nn.modules.conv.Concat	[1]
13		-1	2	1182720	ultralytics.nn.modules.block.C3	[768, 384, 2, False]
14		-1	1	74112	ultralytics.nn.modules.conv.Conv	[384, 192, 1, 1]
15		-1	1	0	torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
16	[-1, 4]	1	1	0	ultralytics.nn.modules.conv.Concat	[1]
17		-1	2	296448	ultralytics.nn.modules.block.C3	[384, 192, 2, False]
18		-1	1	332160	ultralytics.nn.modules.conv.Conv	[192, 192, 3, 2]
19	[-1, 14]	1	1	0	ultralytics.nn.modules.conv.Concat	[1]
20		-1	2	1035264	ultralytics.nn.modules.block.C3	[384, 384, 2, False]
21		-1	1	1327872	ultralytics.nn.modules.conv.Conv	[384, 384, 3, 2]
22	[-1, 10]	1	1	0	ultralytics.nn.modules.conv.Concat	[1]
23		-1	2	4134912	ultralytics.nn.modules.block.C3	[768, 768, 2, False]
24	[17, 20, 23]	1	1	4218643	ultralytics.nn.modules.head.Detect	[1, [192, 384, 768]]

YOLOv5m summary: 339 layers, 25,065,715 parameters, 25,065,699 gradients, 64.4 GFLOPs

Transferred 553/559 items from pretrained weights

TensorBoard: Start with 'tensorboard --logdir runs/detect/car_detection_10', view at <http://localhost:6006/>

Freezing layer 'model.24.dfl.conv.weight'

AMP: running Automatic Mixed Precision (AMP) checks...

AMP: checks passed ✅

```

train: Scanning /blue/eel5934/r.anumula/project-3-graduate-RaviTejaAnumula/car_detection_dataset/train/label
s.cache... 961 images, 625 backgrounds, 0 corrupt: 100%|██████████| 961/961 [00:00<?, ?it/s]
/blue/eel5934/r.anumula/.conda/envs/art/lib/python3.12/site-packages/torch/utils/data/dataloader.py:617: User
Warning: This DataLoader will create 8 worker processes in total. Our suggested max number of worker in curre
nt system is 2, which is smaller than what this DataLoader is going to create. Please be aware that excessive
worker creation might get DataLoader running slow or even freeze, lower the worker number to avoid potential
slowness/freeze if necessary.
  warnings.warn(
val: Scanning /blue/eel5934/r.anumula/project-3-graduate-RaviTejaAnumula/car_detection_dataset/val/labels.cac
he... 362 images, 212 backgrounds, 0 corrupt: 100%|██████████| 362/362 [00:00<?, ?it/s]
/blue/eel5934/r.anumula/.conda/envs/art/lib/python3.12/site-packages/torch/utils/data/dataloader.py:617: User
Warning: This DataLoader will create 16 worker processes in total. Our suggested max number of worker in curr
ent system is 2, which is smaller than what this DataLoader is going to create. Please be aware that excessiv
e worker creation might get DataLoader running slow or even freeze, lower the worker number to avoid potentia
l slowness/freeze if necessary.
  warnings.warn(

```

Plotting labels to runs/detect/car_detection_10/labels.jpg...

optimizer: 'optimizer=auto' found, ignoring 'lr=0.01' and 'momentum=0.937' and determining best 'optimizer', 'lr' and 'momentum' automatically...

optimizer: AdamW(lr=0.002, momentum=0.9) with parameter groups 91 weight(decay=0.0), 98 weight(decay=0.0005), 97 bias(decay=0.0)

TensorBoard: model graph visualization added 

Image sizes 640 train, 640 val

Using 8 dataloader workers

Logging results to runs/detect/car_detection_10

Starting training for 10 epochs...

Closing dataloader mosaic

```

/blue/eel5934/r.anumula/.conda/envs/art/lib/python3.12/site-packages/torch/utils/data/dataloader.py:617: User
Warning: This DataLoader will create 8 worker processes in total. Our suggested max number of worker in curre
nt system is 2, which is smaller than what this DataLoader is going to create. Please be aware that excessive
worker creation might get DataLoader running slow or even freeze, lower the worker number to avoid potential
slowness/freeze if necessary.
  warnings.warn(

```

Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size
1/10	7.07G	1.55	3.891	1.466	0	640: 100% ██████████ 61/61 [00:27<00:00, 2.23it/s]
Class Images Instances Box(P R mAP50 mAP50-95): 100% ██████████ 1 2/12 [00:08<00:00, 1.49it/s]						
	all	362	241	0.111	0.83	0.117 0.0699

Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size
2/10	6.99G	1.581	1.308	1.607	0	640: 100% ██████████ 61/61 [00:08<00:00, 6.92it/s]
Class Images Instances Box(P R mAP50 mAP50-95): 100% ██████████ 1 2/12 [00:01<00:00, 9.67it/s]						
	all	362	241	0.897	0.859	0.928 0.494

Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size
3/10	7.13G	1.593	1.145	1.694	1	640: 100% ██████████ 61/61 [00:08<00:00, 7.03it/s]
Class Images Instances Box(P R mAP50 mAP50-95): 100% ██████████ 1 2/12 [00:01<00:00, 10.48it/s]						
	all	362	241	0.894	0.838	0.92 0.525

Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size
4/10	7.03G	1.485	1.066	1.609	0	640: 100% ██████████ 61/61 [00:08<00:00, 7.10it/s]
Class Images Instances Box(P R mAP50 mAP50-95): 100% ██████████ 1 2/12 [00:01<00:00, 10.14it/s]						
	all	362	241	0.178	0.765	0.173 0.0991

Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size
5/10	7.13G	1.395	0.8885	1.52	1	640: 100% ██████████ 61/61 [00:08<00:00, 7.03it/s]
Class Images Instances Box(P R mAP50 mAP50-95): 100% ██████████ 1 2/12 [00:01<00:00, 10.24it/s]						
	all	362	241	0.888	0.888	0.945 0.562

Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size
-------	---------	----------	----------	----------	-----------	------

6/10	7.1G	1.433	0.8864	1.531	0	640: 100%	<div></div>	61/61 [00:08<00:00, 7.18it/s]
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% <div></div> 1		
2/12 [00:01<00:00, 10.04it/s]	all	362	241	0.948	0.903	0.972	0.635	
Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size		
7/10	7.08G	1.373	0.7914	1.512	0	640: 100%	<div></div>	61/61 [00:08<00:00, 7.21it/s]
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% <div></div> 1		
2/12 [00:01<00:00, 10.07it/s]	all	362	241	0.932	0.925	0.974	0.612	
Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size		
8/10	7G	1.305	0.7805	1.391	2	640: 100%	<div></div>	61/61 [00:08<00:00, 6.81it/s]
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% <div></div> 1		
2/12 [00:01<00:00, 10.82it/s]	all	362	241	0.948	0.975	0.988	0.643	
Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size		
9/10	7.13G	1.293	0.7106	1.405	0	640: 100%	<div></div>	61/61 [00:08<00:00, 7.16it/s]
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% <div></div> 1		
2/12 [00:01<00:00, 11.28it/s]	all	362	241	0.95	0.938	0.986	0.665	
Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size		
10/10	7.12G	1.248	0.6657	1.399	1	640: 100%	<div></div>	61/61 [00:08<00:00, 7.23it/s]
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% <div></div> 1		
2/12 [00:01<00:00, 10.29it/s]	all	362	241	0.971	0.954	0.991	0.657	
10 epochs completed in 0.039 hours.								
Optimizer stripped from runs/detect/car_detection_10/weights/last.pt, 50.5MB								
Optimizer stripped from runs/detect/car_detection_10/weights/best.pt, 50.5MB								
Validating runs/detect/car_detection_10/weights/best.pt...								
Ultralytics 8.3.39 Python-3.12.7 torch-2.5.1+cu124 CUDA:0 (NVIDIA A100-SXM4-80GB, 81038MiB)								
YOLOv5m summary (fused): 248 layers, 25,045,795 parameters, 0 gradients, 64.0 GFLOPs								
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% <div></div> 1		
2/12 [00:01<00:00, 8.33it/s]								

all 362 241 0.95 0.938 0.986 0.664
Speed: 0.1ms preprocess, 1.0ms inference, 0.0ms loss, 0.6ms postprocess per image

Results saved to **runs/detect/car_detection_10**

PRO TIP 🟡 Replace 'model=yolov5m.pt' with new 'model=yolov5mu.pt'.

YOLOv5 'u' models are trained with <https://github.com/ultralytics/ultralytics> and feature improved performance vs standard YOLOv5 models trained with <https://github.com/ultralytics/yolov5>.

New <https://pypi.org/project/ultralytics/8.3.41> available 🤗 Update with 'pip install -U ultralytics'

Ultralytics 8.3.39 🚀 Python-3.12.7 torch-2.5.1+cu124 CUDA:0 (NVIDIA A100-SXM4-80GB, 81038MiB)

engine/trainer: task=detect, mode=train, model=yolov5m.pt, data=data.yaml, epochs=20, time=None, patience=100, batch=16, imgsz=640, save=True, save_period=-1, cache=False, device=None, workers=8, project=None, name=car_detection_20, exist_ok=False, pretrained=True, optimizer=auto, verbose=True, seed=0, deterministic=True, single_cls=False, rect=False, cos_lr=False, close_mosaic=10, resume=False, amp=True, fraction=1.0, profile=False, freeze=None, multi_scale=False, overlap_mask=True, mask_ratio=4, dropout=0.0, val=True, split=val, save_json=False, save_hybrid=False, conf=None, iou=0.7, max_det=300, half=False, dnn=False, plots=True, source=None, vid_stride=1, stream_buffer=False, visualize=False, augment=False, agnostic_nms=False, classes=None, retina_masks=False, embed=None, show=False, save_frames=False, save_txt=False, save_conf=False, save_crop=False, show_labels=True, show_conf=True, show_boxes=True, line_width=None, format=torchscript, keras=False, optimize=False, int8=False, dynamic=False, simplify=True, opset=None, workspace=None, nms=False, lr0=0.01, lrf=0.01, momentum=0.937, weight_decay=0.0005, warmup_epochs=3.0, warmup_momentum=0.8, warmup_bias_lr=0.1, box=7.5, cls=0.5, dfl=1.5, pose=12.0, kobj=1.0, nbs=64, hsv_h=0.015, hsv_s=0.7, hsv_v=0.4, degrees=0.0, translate=0.1, scale=0.5, shear=0.0, perspective=0.0, flipud=0.0, fliplr=0.5, bgr=0.0, mosaic=1.0, mixup=0.0, copy_paste=0.0, copy_paste_mode=flip, auto_augment=randaugment, erasing=0.4, crop_fraction=1.0, cfg=None, tracker=botsort.yaml, save_dir=runs/detect/car_detection_20

Overriding model.yaml nc=80 with nc=1

		from	n	params	module	arguments
0		-1	1	5280	ultralytics.nn.modules.conv.Conv	[3, 48, 6, 2, 2]
1		-1	1	41664	ultralytics.nn.modules.conv.Conv	[48, 96, 3, 2]
2		-1	2	65280	ultralytics.nn.modules.block.C3	[96, 96, 2]
3		-1	1	166272	ultralytics.nn.modules.conv.Conv	[96, 192, 3, 2]
4		-1	4	444672	ultralytics.nn.modules.block.C3	[192, 192, 4]
5		-1	1	664320	ultralytics.nn.modules.conv.Conv	[192, 384, 3, 2]
6		-1	6	2512896	ultralytics.nn.modules.block.C3	[384, 384, 6]
7		-1	1	2655744	ultralytics.nn.modules.conv.Conv	[384, 768, 3, 2]
8		-1	2	4134912	ultralytics.nn.modules.block.C3	[768, 768, 2]
9		-1	1	1476864	ultralytics.nn.modules.block.SPPF	[768, 768, 5]
10		-1	1	295680	ultralytics.nn.modules.conv.Conv	[768, 384, 1, 1]
11		-1	1	0	torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
12		[-1, 6]	1	0	ultralytics.nn.modules.conv.Concat	[1]
13		-1	2	1182720	ultralytics.nn.modules.block.C3	[768, 384, 2, False]
14		-1	1	74112	ultralytics.nn.modules.conv.Conv	[384, 192, 1, 1]
15		-1	1	0	torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
16		[-1, 4]	1	0	ultralytics.nn.modules.conv.Concat	[1]
17		-1	2	296448	ultralytics.nn.modules.block.C3	[384, 192, 2, False]
18		-1	1	332160	ultralytics.nn.modules.conv.Conv	[192, 192, 3, 2]
19		[-1, 14]	1	0	ultralytics.nn.modules.conv.Concat	[1]
20		-1	2	1035264	ultralytics.nn.modules.block.C3	[384, 384, 2, False]
21		-1	1	1327872	ultralytics.nn.modules.conv.Conv	[384, 384, 3, 2]
22		[-1, 10]	1	0	ultralytics.nn.modules.conv.Concat	[1]
23		-1	2	4134912	ultralytics.nn.modules.block.C3	[768, 768, 2, False]
24		[17, 20, 23]	1	4218643	ultralytics.nn.modules.head.Detect	[1, [192, 384, 768]]

YOLOv5m summary: 339 layers, 25,065,715 parameters, 25,065,699 gradients, 64.4 GFLOPs

Transferred 553/559 items from pretrained weights

TensorBoard: Start with 'tensorboard --logdir runs/detect/car_detection_20', view at <http://localhost:6006/>

Freezing layer 'model.24.dfl.conv.weight'

AMP: running Automatic Mixed Precision (AMP) checks...

AMP: checks passed ✅

```

train: Scanning /blue/eel5934/r.anumula/project-3-graduate-RaviTejaAnumula/car_detection_dataset/train/label
s.cache... 961 images, 625 backgrounds, 0 corrupt: 100%|██████████| 961/961 [00:00<?, ?it/s]
/blue/eel5934/r.anumula/.conda/envs/art/lib/python3.12/site-packages/torch/utils/data/dataloader.py:617: User
Warning: This DataLoader will create 8 worker processes in total. Our suggested max number of worker in curre
nt system is 2, which is smaller than what this DataLoader is going to create. Please be aware that excessive
worker creation might get DataLoader running slow or even freeze, lower the worker number to avoid potential
slowness/freeze if necessary.
  warnings.warn(
val: Scanning /blue/eel5934/r.anumula/project-3-graduate-RaviTejaAnumula/car_detection_dataset/val/labels.cac
he... 362 images, 212 backgrounds, 0 corrupt: 100%|██████████| 362/362 [00:00<?, ?it/s]
/blue/eel5934/r.anumula/.conda/envs/art/lib/python3.12/site-packages/torch/utils/data/dataloader.py:617: User
Warning: This DataLoader will create 16 worker processes in total. Our suggested max number of worker in curr
ent system is 2, which is smaller than what this DataLoader is going to create. Please be aware that excessiv
e worker creation might get DataLoader running slow or even freeze, lower the worker number to avoid potentia
l slowness/freeze if necessary.
  warnings.warn(
Plotting labels to runs/detect/car_detection_20/labels.jpg...
optimizer: 'optimizer=auto' found, ignoring 'lr=0.01' and 'momentum=0.937' and determining best 'optimizer',
'lr' and 'momentum' automatically...
optimizer: AdamW(lr=0.002, momentum=0.9) with parameter groups 91 weight(decay=0.0), 98 weight(decay=0.0005),
97 bias(decay=0.0)
TensorBoard: model graph visualization added ✅
Image sizes 640 train, 640 val
Using 8 dataloader workers
Logging results to runs/detect/car_detection_20
Starting training for 20 epochs...

```

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size		
1/20	7.12G	1.543	1.915	1.402	0	640:	100% ██████████	61/61 [00:09<00:00, 6.30it/s]
		Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 1
2/12	[00:01<00:00, 9.58it/s]							
		all	362	241	0.846	0.89	0.924	0.554
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size		
2/20	7.03G	1.424	1.087	1.416	0	640:	100% ██████████	61/61 [00:09<00:00, 6.70it/s]
		Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 1
2/12	[00:01<00:00, 9.55it/s]							
		all	362	241	0.112	0.606	0.0947	0.0547
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size		
3/20	7.14G	1.455	1.031	1.476	0	640:	100% ██████████	61/61 [00:08<00:00, 6.80it/s]
		Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 1
2/12	[00:01<00:00, 10.36it/s]							
		all	362	241	0.0382	0.776	0.0357	0.0191
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size		
4/20	6.99G	1.426	0.9197	1.474	2	640:	100% ██████████	61/61 [00:08<00:00, 6.87it/s]
		Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 1
2/12	[00:01<00:00, 9.43it/s]							
		all	362	241	0.254	0.232	0.0931	0.051
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size		
5/20	7.14G	1.371	0.826	1.426	3	640:	100% ██████████	61/61 [00:09<00:00, 6.65it/s]
		Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 1
2/12	[00:01<00:00, 10.79it/s]							
		all	362	241	0.854	0.874	0.941	0.579
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size		
6/20	7.11G	1.365	0.7784	1.408	0	640:	100% ██████████	61/61 [00:08<00:00, 6.86it/s]
		Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 1
2/12	[00:01<00:00, 9.24it/s]							
		all	362	241	0.933	0.922	0.963	0.554
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size		

7/20	7.09G	1.369	0.7968	1.395	1	640: 100%		61/61 [00:08<00:00, 6.82it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100%	1
2/12	[00:01<00:00, 10.22it/s]							
	all	362	241	0.936	0.964	0.987	0.632	
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size		
8/20	7.03G	1.282	0.7111	1.342	0	640: 100%		61/61 [00:08<00:00, 6.82it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100%	1
2/12	[00:01<00:00, 9.25it/s]							
	all	362	241	0.975	0.952	0.989	0.634	
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size		
9/20	7.14G	1.269	0.7029	1.31	0	640: 100%		61/61 [00:08<00:00, 6.83it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100%	1
2/12	[00:01<00:00, 9.81it/s]							
	all	362	241	0.963	0.992	0.99	0.653	
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size		
10/20	7.1G	1.208	0.6465	1.286	0	640: 100%		61/61 [00:09<00:00, 6.65it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100%	1
2/12	[00:01<00:00, 10.62it/s]							
	all	362	241	0.955	0.983	0.99	0.654	
Closing dataloader mosaic								
/blue/eel5934/r.anumula/.conda/envs/art/lib/python3.12/site-packages/torch/utils/data/dataloader.py:617: User Warning: This DataLoader will create 8 worker processes in total. Our suggested max number of worker in current system is 2, which is smaller than what this DataLoader is going to create. Please be aware that excessive worker creation might get DataLoader running slow or even freeze, lower the worker number to avoid potential slowness/freeze if necessary.								
warnings.warn(
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size		
11/20	7.1G	1.327	0.7113	1.415	1	640: 100%		61/61 [00:10<00:00, 5.99it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100%	1
2/12	[00:01<00:00, 11.18it/s]							
	all	362	241	0.95	0.948	0.988	0.666	
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size		
12/20	7G	1.263	0.7012	1.372	0	640: 100%		61/61 [00:08<00:00, 7.18it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100%	1
2/12	[00:01<00:00, 9.62it/s]							
	all	362	241	0.926	0.931	0.976	0.642	
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size		
13/20	7.14G	1.251	0.6748	1.38	2	640: 100%		61/61 [00:08<00:00, 7.18it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100%	1
2/12	[00:01<00:00, 10.26it/s]							
	all	362	241	0.971	0.975	0.991	0.657	
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size		
14/20	7.09G	1.237	0.6221	1.394	0	640: 100%		61/61 [00:08<00:00, 7.15it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100%	1
2/12	[00:01<00:00, 10.29it/s]							
	all	362	241	0.959	0.946	0.988	0.667	
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size		
15/20	7.08G	1.274	0.6249	1.38	0	640: 100%		61/61 [00:08<00:00, 7.09it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100%	1
2/12	[00:01<00:00, 9.52it/s]							
	all	362	241	0.975	0.979	0.993	0.674	
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size		

Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size	
16/20	7.19G	1.223	0.5891	1.378	0	640: 100% ██████████ 61/61 [00:08<00:00, 7.19it/s]	
		Class	Images	Instances	Box(P	R	mAP50 mAP50-95): 100% ██████████ 1
2/12 [00:01<00:00, 10.01it/s]		all	362	241	0.96	0.984	0.99 0.671
17/20	7.14G	1.21	0.5566	1.337	0	640: 100% ██████████ 61/61 [00:08<00:00, 7.17it/s]	
		Class	Images	Instances	Box(P	R	mAP50 mAP50-95): 100% ██████████ 1
2/12 [00:01<00:00, 9.81it/s]		all	362	241	0.971	0.986	0.992 0.689
18/20	7.09G	1.176	0.5698	1.331	1	640: 100% ██████████ 61/61 [00:08<00:00, 7.09it/s]	
		Class	Images	Instances	Box(P	R	mAP50 mAP50-95): 100% ██████████ 1
2/12 [00:01<00:00, 10.82it/s]		all	362	241	0.96	0.992	0.992 0.687
19/20	7.1G	1.147	0.528	1.348	1	640: 100% ██████████ 61/61 [00:08<00:00, 7.07it/s]	
		Class	Images	Instances	Box(P	R	mAP50 mAP50-95): 100% ██████████ 1
2/12 [00:01<00:00, 11.47it/s]		all	362	241	0.971	0.987	0.993 0.701
20/20	7.02G	1.136	0.5283	1.287	0	640: 100% ██████████ 61/61 [00:08<00:00, 7.17it/s]	
		Class	Images	Instances	Box(P	R	mAP50 mAP50-95): 100% ██████████ 1
2/12 [00:01<00:00, 10.43it/s]		all	362	241	0.972	0.996	0.993 0.703

20 epochs completed in 0.095 hours.
 Optimizer stripped from runs/detect/car_detection_20/weights/last.pt, 50.5MB
 Optimizer stripped from runs/detect/car_detection_20/weights/best.pt, 50.5MB

Validating runs/detect/car_detection_20/weights/best.pt...
 Ultralytics 8.3.39 Python-3.12.7 torch-2.5.1+cu124 CUDA:0 (NVIDIA A100-SXM4-80GB, 81038MiB)
 YOLOv5m summary (fused): 248 layers, 25,045,795 parameters, 0 gradients, 64.0 GFLOPs

Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100%
2/12 [00:01<00:00, 8.90it/s]	all	362	241	0.972	0.996	0.993 0.703

Speed: 0.1ms preprocess, 1.0ms inference, 0.0ms loss, 0.6ms postprocess per image
 Results saved to runs/detect/car_detection_20

In [11]: # If a best model was found, save its information

```
if best_model_info:
    model_path = best_model_info['model_path']
    results_csv = best_model_info['results_csv']
    with open(pickle_file, 'wb') as f:
        pickle.dump(best_model_info, f)
```

In [14]: # If best model info is available, plot Learning curves
 # Assuming best_model_info is defined and df is loaded correctly

```
if best_model_info:
    df = pd.read_csv(best_model_info["results_csv"])
    metrics = ['train/box_loss', 'train/cls_loss', 'metrics/precision(B)', 'metrics/recall(B)']
    # Create subplots
    fig, axs = plt.subplots(2, 2, figsize=(15, 15))
    axs = axs.ravel() # Flatten the 2D array to 1D for easier iteration
    # Loop through each metric and plot it
    for i, metric in enumerate(metrics):
        if metric in df.columns:
            axs[i].plot(df['epoch'], df[metric], label=metric)
            axs[i].set_title(f'{metric} vs Epoch')
            axs[i].set_xlabel('Epoch')
            axs[i].set_ylabel('Value')
            axs[i].legend(loc='upper right')
        else:
            axs[i].axis('off') # Hide the axis if the metric is missing
```

```
plt.tight_layout()
# Save figure as PNG
plt.savefig('learning-curves.png')
print("Learning curves saved as 'learning-curves.png'")
# Show the plot
plt.show()
print(f"Best mAP achieved: {best_map}")
```

Learning curves saved as 'learning-curves.png'
 <Figure size 1500x1500 with 4 Axes>
 Best mAP achieved: 0.97215

3

To assess how well the model performs without relying solely on labels, you can involve a small group of people for additional validation. Select a small set of test images—say, 50 or 100—and show these images to a few individuals, asking them to mark areas where they see cars. Afterward, compare their annotations with the predictions made by your model. This will give you a clearer picture of how accurate your model's predictions are in real-world scenarios.

For images without cars, you can train the model to recognize empty scenes as well. When testing, if the model doesn't detect any cars in an image, you can verify this by reviewing the image manually or asking your human helpers to confirm.

This approach provides a practical way to gauge your model's effectiveness, especially in cases where some images contain cars and others do not. It acts as a reality check, using human judgment to evaluate how well the model is performing.

```
In [2]: df = pd.read_csv('./car_detection_dataset/train_bounding_boxes.csv')
image_files = sorted([l for l in os.listdir('./car_detection_dataset/training_images')
                      if l.lower().endswith(('.png', '.jpg', '.jpeg'))
                      and os.path.isfile(os.path.join('./car_detection_dataset/training_images', l))])

train_images, val_images = train_test_split(image_files, test_size=0.2, random_state=42)

os.makedirs('./car_detection_dataset/train_images_with_and_without_cars/images', exist_ok=True)
os.makedirs('./car_detection_dataset/train_images_with_and_without_cars/labels', exist_ok=True)
os.makedirs('./car_detection_dataset/val_images_with_and_without_cars/images', exist_ok=True)
os.makedirs('./car_detection_dataset/val_images_with_and_without_cars/labels', exist_ok=True)

In [3]: # Process and move training images
for images in train_images:
    source = os.path.join('./car_detection_dataset/training_images', images)

    # Copy image to the destination folder
    destination = os.path.join('./car_detection_dataset/train_images_with_and_without_cars', images)
    shutil.copy(source, destination)

    # Retrieve bounding box data for the current image
    image_data = df[df['image'] == images]
    label_path = os.path.join('./car_detection_dataset/train_images_with_and_without_cars/labels', f'{Path(i

    # Open the label file and write the bounding box annotations
    with open(label_path, 'w') as lp:
        if not image_data.empty:
            image = cv2.imread(source)
            for _, row in image_data.iterrows():
                # Normalize the bounding box coordinates
                x = (row['xmin'] + row['xmax']) / 2 / image.shape[1]
                y = (row['ymin'] + row['ymax']) / 2 / image.shape[0]
                w = (row['xmax'] - row['xmin']) / image.shape[1]
                h = (row['ymax'] - row['ymin']) / image.shape[0]

                # Write the annotation (label) in YOLO format
                lp.write(f'0 {x} {y} {w} {h}\n')
```

```
In [4]: # Process and move validation images
for images in val_images:
    source = os.path.join('./car_detection_dataset/training_images', images)
    # Copy image to the destination folder
    destination = os.path.join('./car_detection_dataset/val_images_with_and_without_cars', images)
    shutil.copy(source, destination)

    # Retrieve bounding box data for the current image
    image_data = df[df['image'] == images]
    label_path = os.path.join('./car_detection_dataset/val_images_with_and_without_cars/labels', f'{Path(ima

    # Open the label file and write the bounding box annotations
    with open(label_path, 'w') as lp:
        if not image_data.empty:
            image = cv2.imread(source)
            for _, row in image_data.iterrows():
                # Normalize the bounding box coordinates
                x = (row['xmin'] + row['xmax']) / 2 / image.shape[1]
                y = (row['ymin'] + row['ymax']) / 2 / image.shape[0]
                w = (row['xmax'] - row['xmin']) / image.shape[1]
                h = (row['ymax'] - row['ymin']) / image.shape[0]

                # Write the annotation (label) in YOLO format
                lp.write(f'0 {x} {y} {w} {h}\n')
```

```
In [5]: data_yaml = {
    'train': './car_detection_dataset/train_images_with_and_without_cars/images',
    'val': './car_detection_dataset/val_images_with_and_without_cars/images',
    'nc': 1, # Number of classes (in this case, only 'car')
    'names': ['car']
}

# Save the YAML configuration to a file
import yaml
with open('dataset_with_and_without_cars.yaml', 'w') as f:
    yaml.dump(data_yaml, f)
```

```
In [6]: maxWidth, maxHeight = 0, 0
for name in os.listdir('./car_detection_dataset/training_images/'):
    imagePath = os.path.join('./car_detection_dataset/training_images/', name)
    # Read the image
    image = cv2.imread(imagePath)
    if image is not None: # Ensure the image was read successfully
        height = image.shape[0]
        width = image.shape[1] # Get the dimensions
        # Update max width and height if necessary
        if width > maxWidth:
            maxWidth = width
        if height > maxHeight:
            maxHeight = height

    # Output the maximum width and height
    print(f"Max Width: {maxWidth}, Max Height: {maxHeight}")
```

Max Width: 676, Max Height: 380

```
In [7]: model = YOLO('yolov5m.pt')
imgsz = max(32 * round(max(maxWidth, maxHeight) / 32), 640)
model_path = './runs/detect/car_detection_20/weights/best.pt'
```

PRO TIP 💡 Replace 'model=yolov5m.pt' with new 'model=yolov5mu.pt'.
YOLOv5 'u' models are trained with <https://github.com/ultralytics/ultralytics> and feature improved performance vs standard YOLOv5 models trained with <https://github.com/ultralytics/yolov5>.

```
In [8]: def evaluate_predictions(pred_boxes, true_boxes, image, iou_threshold=0.8, buffer=4):
    cars_identified = 0
    total_cars = len(true_boxes)

    img_h = image.shape[0]
```

```

img_w = image.shape[1]

for true_box in true_boxes:
    x, y, w, h = true_box
    x1 = int((x - w/2) * img_w)
    y1 = int((y - h/2) * img_h)
    x2 = int((x + w/2) * img_w)
    y2 = int((y + h/2) * img_h)
    true_box_coords = [x1, y1, x2, y2]

    b2 = [x1 - buffer, y1 - buffer, x2 + buffer, y2 + buffer]
    for pred_box in pred_boxes:
        pred_x1, pred_y1, pred_x2, pred_y2 = pred_box

        # Calculate intersection area
        inter_x1 = max(pred_x1, b2[0])
        inter_y1 = max(pred_y1, b2[1])
        inter_x2 = min(pred_x2, b2[2])
        inter_y2 = min(pred_y2, b2[3])

        inter_area = max(0, inter_x2 - inter_x1) * max(0, inter_y2 - inter_y1)
        pred_area = (pred_x2 - pred_x1) * (pred_y2 - pred_y1)
        true_area = (b2[2] - b2[0]) * (b2[3] - b2[1])

        iou = inter_area / float(pred_area + true_area - inter_area)

        if iou > iou_threshold:
            cars_identified += 1
            break
return cars_identified, total_cars

```

```

In [10]: def load_true_boxes(label_path):
    """Load ground truth bounding boxes from a label file."""
    try:
        with open(label_path, 'r') as f:
            return [list(map(float, line.strip().split()[1:])) for line in f]
    except FileNotFoundError:
        return []

def evaluate_and_accumulate(image_path, label_path, model, iou_threshold=0.8, buffer=4):
    """Evaluate predictions on a single image and accumulate identified and total car counts."""
    image = cv2.imread(image_path)
    predictions = model(image_path, verbose=False)
    pred_boxes = predictions[0].boxes.xyxy.cpu().numpy()

    true_boxes = load_true_boxes(label_path)

    cars_identified, cars_in_image = evaluate_predictions(pred_boxes, true_boxes, image, iou_threshold, buff

    return cars_identified, cars_in_image

def calculate_accuracy(total_identified, total_actual):
    """Calculate accuracy based on total identified and total actual cars."""
    if total_actual > 0:
        return total_identified / total_actual
    else:
        return 0.0

def main(image_dir, label_dir, model_path, iou_threshold=0.8, buffer=4):
    model = YOLO(model_path)

    total_cars_identified = 0
    total_cars = 0

    for image_file in os.listdir(image_dir):
        if not image_file.lower().endswith(('.png', '.jpg', '.jpeg')):
            continue

        image_path = os.path.join(image_dir, image_file)
        label_path = os.path.join(label_dir, os.path.splitext(image_file)[0] + '.txt')

```

```
cars_identified, cars_in_image = evaluate_and_accumulate(image_path, label_path, model, iou_threshol

total_cars_identified += cars_identified
total_cars += cars_in_image

accuracy = calculate_accuracy(total_cars_identified, total_cars)

print(f"Cars identified: {total_cars_identified}")
print(f"Total cars: {total_cars}")
print(f"Accuracy: {accuracy:.2f}")

# Define directories and model path
image_dir = './Make sense/'
label_dir = './Make sense labels/'
model_path = './runs/detect/car_detection_20/weights/best.pt'

# Run the evaluation
main(image_dir, label_dir, model_path)
```

Cars identified: 1
Total cars: 12
Accuracy: 0.08

In []: